B20 review brief: learning to research (W18–W20, E3)

The B20 exit is "better-improver claim supported or explicitly unestablished". The builder reports UNESTABLISHED. This brief maps what a reviewer has to check.

Reproduce

```
export GENESIS_TOOLCHAIN=/srv/toolchain PYTHONPATH=src PYTHONPYCACHEPREFIX=$(mkte
python3 -m unittest tests_formal.test_learning
python3 scripts/b20_learning.py /tmp/b20                  # about 40 min on 2 cor
python3 scripts/b20_e3_hard.py /tmp/b20 311 312 313 314    # needs /tmp/b20/train
python3 scripts/b20_hard_followup.py /tmp/b20
python3 scripts/b20_curriculum_control.py /tmp/b20
```

E runs with CPU limits. Re-runs on a loaded host can differ: see the concurrency caveat in the evidence README.

Decisions for the reviewer

#	Decision	Evidence
1	Are the deciders sound ground truth? Boolean algebra is evaluated over {0,1} (the variety is generated by the 2-element algebra), and groups by free-group reduction.	`equational.py`; `Deciders` tests
2	Does a lemma enter a library only with a proof, and are successor proofs Lean-checked down to the axioms?	`research.run`, `lean_chain`; `LemmaChains` tests; `lean/replays.json`
3	Is the policy comparison fair? Equal CPU budgets, fresh seeds, the learned policy frozen before evaluation, and evaluation targets never visible to research.	`b20_learning.py`; `setup.json`, `train/policy.json`
4	Were the post-hoc additions handled honestly? These are the hard-system E3, the follow-up, non-strict	the script docstrings; README

	expansion and the curriculum control.	
5	Is screening adequate to rule out near-duplicate answers?	`memory.screen` ; the exclusions in `w20/rows.json` and `curriculum_control.json`

Known limits

- **Domain.** Research has room only on Wolfram's axiom, where single runs are noisy.
- **E3.** Learned ≈ fixed; nothing here supports a better-improver claim.
- **Macros.** They do not beat plain research.